

Introduction to Monte-Carlo Methods

Rohan L. Fernando

May 2015

Mean and Variance of Truncated Normal

Suppose $Y \sim N(\mu_Y, V_Y)$.

The mean and variance of Y given truncation selection are:

$$E(Y|Y > t) = \mu_Y + V_Y^{1/2}i$$

where

$$i = \frac{f(s)}{p}$$

$f(s)$ is the standard normal density function

$$s = \frac{t - \mu_Y}{V_Y^{1/2}}$$

$$p = \Pr(Y > t)$$

$$\text{Var}(Y|Y > t) = V_Y[1 - i(i - s)]$$

Proof:

Start with mean and variance for a standard normal variable given truncation selection.

Let $Z \sim N(0, 1)$.

The density function of Z is:

$$f(z) = \sqrt{\frac{1}{2\pi}} e^{-\frac{1}{2}z^2}$$

The density function for Z given truncation selection is

$$f(z|z > s) = f(z)/p$$

From the definition of the mean:

$$\begin{aligned}
 E(Z|Z > s) &= \frac{1}{p} \int_s^\infty zf(z)dz \\
 &= \frac{1}{p} [-f(z)]_s^\infty \\
 &= \frac{f(s)}{p} \\
 &= i
 \end{aligned}$$

because the first derivative of $f(z)$ with respect to z is:

$$\begin{aligned}
 \frac{d}{dz}f(z) &= \sqrt{\frac{1}{2\pi}} e^{-\frac{1}{2}z^2} (-z) \\
 &= -zf(z)
 \end{aligned}$$

Now, to compute the variance of Z given selection, consider the following identity:

$$\begin{aligned}
 \frac{d}{dz}zf(z) &= f(z) + z\frac{d}{dz}f(z) \\
 &= f(z) - z^2f(z)
 \end{aligned}$$

Integrating both sides from s to ∞ gives

$$zf(z)]_s^\infty = \int_s^\infty f(z)dz - \int_s^\infty z^2f(z)dz$$

Upon rearranging this gives:

$$\begin{aligned}
 \int_s^\infty z^2f(z)dz &= \int_s^\infty f(z)dz - zf(z)]_s^\infty \\
 \frac{1}{p} \int_s^\infty z^2f(z)dz &= \frac{1}{p} \int_s^\infty f(z)dz + \frac{f(s)}{p}s \\
 &= 1 + is
 \end{aligned}$$

So,

$$\begin{aligned}
 Var(Z|Z > s) &= E(Z^2|Z > s) - [E(Z|Z > s)]^2 \\
 &= 1 + is - i^2 \\
 &= 1 - i(i - s)
 \end{aligned}$$

Results for Y

Results for Y follow from the fact that

$$\mu_Y + V_Y^{1/2}Z \sim N(\mu_Y, V_Y)$$

So, let

$$Y = \mu_Y + V_Y^{1/2}Z,$$

Then, the condition

$$Y > t$$

is equivalent to

$$\begin{aligned}\mu_Y + V_Y^{1/2}Z &> t \\ V_Y^{1/2}Z &> t - \mu_Y \\ Z &> \frac{t - \mu_Y}{V_Y^{1/2}} \\ Z &> s\end{aligned}$$

Then,

$$\begin{aligned}E(Y|Y > t) &= E(\mu_Y + V_Y^{1/2}Z|Z > s) \\ &= \mu_Y + V_Y^{1/2}i,\end{aligned}$$

and

$$\begin{aligned}Var(Y|Y > t) &= Var(\mu_Y + V_Y^{1/2}Z|Z > s) \\ &= V_Y[1 - i(i - s)]\end{aligned}$$

Numerical Example

```
In [39]: μ = 10
σ = 10
t = 15
s = (t-μ)/σ
d = Normal(0.0,1.0)
i = pdf(d,s)/(1-cdf(d,s))
meanTruncatedNormal = μ + σ*i
variTruncatedNormal = σ*σ*(1 - i*(i-s))
@printf "mean      = %8.2f  \n" meanTruncatedNormal
@printf "variance = %8.2f  \n" variTruncatedNormal

mean      =      21.41
variance =      26.85
```

Monte-Carlo Approach:

```
In [43]: using Distributions
μ = 10
σ = 10
z = rand(Normal(μ,σ),10000);

In [56]: mcmcMean = mean(z[z.>t])
mcmcVar = var(z[z.>t])
@printf "MC mean      = %8.2f  \n" mcmcMean
@printf "MC variance = %8.2f  \n" mcmcVar

MC mean      =      21.34
MC variance =      25.78
```

Bivariate Normal Example

Let $(Y) \sim N(\mu, \mathbf{V})$

$$\mu = \begin{bmatrix} 10 \\ 20 \end{bmatrix}, \mathbf{V} = \begin{bmatrix} 100 & 50 \\ 50 & 200 \end{bmatrix}$$

```
In [54]: μ = [10.0;20.0]
V = [100.0 50.0
      50.0 200.0]
d = MvNormal(μ,V)
XY = rand(d,10000)'
```

Out[54]: 10000x2 Array{Float64,2}:

```
10.3117    41.2371
 8.49604   30.121
 1.49591    5.04669
 2.0137    21.2858
 8.12043    9.99512
17.9018    16.9568
 1.01726    20.0321
-8.29162    40.2454
14.6496    45.1535
13.9381    12.9118
-0.612875  24.1609
20.5875    15.1366
16.2409    25.9275
⋮
 3.98896    3.67185
13.8927    24.0219
 3.93784    11.8521
 3.83364    4.41762
20.7947    37.1139
 9.11036    15.7678
 4.45919    32.2166
19.5114    21.9018
12.777     29.3537
18.1348    11.6092
 0.640994   14.6436
 3.39195    27.4398
```

```
In [111]: sel = XY[:,1].>10  
xxy= [XY sel]
```

```
Out[111]: 10000x3 Array{Float64,2}:  
 10.3117    41.2371    1.0  
  8.49604   30.121     0.0  
  1.49591    5.04669    0.0  
  2.0137    21.2858    0.0  
  8.12043    9.99512    0.0  
 17.9018    16.9568    1.0  
  1.01726   20.0321    0.0  
 -8.29162   40.2454    0.0  
 14.6496    45.1535    1.0  
 13.9381    12.9118    1.0  
 -0.612875  24.1609    0.0  
 20.5875    15.1366    1.0  
 16.2409    25.9275    1.0  
  ⋮  
  3.98896    3.67185    0.0  
 13.8927    24.0219    1.0  
  3.93784    11.8521    0.0  
  3.83364    4.41762    0.0  
 20.7947    37.1139    1.0  
  9.11036    15.7678    0.0  
  4.45919    32.2166    0.0  
 19.5114    21.9018    1.0  
 12.777     29.3537    1.0  
 18.1348    11.6092    1.0  
  0.640994   14.6436    0.0  
  3.39195    27.4398    0.0
```

```
In [115]: (xxy[:,1][xxy[:,3].==1])
```

```
Out[115]: 18.03854352069298
```

```
In [59]: selY = XY[sel,2]
```

```
Out[59]: 5026-element Array{Float64,1}:  
 41.2371  
 16.9568  
 45.1535  
 12.9118  
 15.1366  
 25.9275  
 17.4284  
 20.6601  
 44.2587  
  7.21451  
 26.9525  
 29.502  
 41.1791  
  ⋮  
 41.4734  
 20.1128  
 33.6962  
 17.7152  
 16.6372  
 48.6728  
 27.0785  
 24.0219  
 37.1139  
 21.9018  
 29.3537  
 11.6092
```

```
In [60]: mean(selY[selY.>30])
```

```
Out[60]: 38.95540792778809
```

```
In [61]: var(selY[selY.>30])
```

```
Out[61]: 52.61527300087836
```

Markov Chain Monte-Carlo Methods

- Often no closed form for $f(\theta|y)$
- Further, even if computing $f(\theta|y)$ is feasible, obtaining $f(\theta_i|y)$ would require integrating over many dimensions
- Thus, in many situations, inferences are made using the empirical posterior constructed by drawing samples from $f(\theta|y)$
- Gibbs sampler is widely used for drawing samples from posteriors

Gibbs Sampler

- Want to draw samples from $f(x_1, x_2, \dots, x_n)$
- Even though it may be possible to compute $f(x_1, x_2, \dots, x_n)$, it is difficult to draw samples directly from $f(x_1, x_2, \dots, x_n)$
- Gibbs:
 - Get valid a starting point \mathbf{x}^0
 - Draw sample \mathbf{x}^t as:

$$\begin{array}{ll} x_1^t & \text{from } f(x_1 | x_2^{t-1}, x_3^{t-1}, \dots, x_n^{t-1}) \\ x_2^t & \text{from } f(x_2 | x_1^t, x_3^{t-1}, \dots, x_n^{t-1}) \\ x_3^t & \text{from } f(x_3 | x_1^t, x_2^t, \dots, x_n^{t-1}) \\ \vdots & \vdots \\ x_n^t & \text{from } f(x_n | x_1^t, x_2^t, \dots, x_{n-1}^t) \end{array}$$
- The sequence $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$ is a Markov chain with stationary distribution $f(x_1, x_2, \dots, x_n)$

Making Inferences from Markov Chain

Can show that samples obtained from a Markov chain can be used to draw inferences from $f(x_1, x_2, \dots, x_n)$ provided the chain is:

- Irreducible: can move from any state i to any other state j
- Positive recurrent: return time to any state has finite expectation
- *Markov Chains*, J. R. Norris (1997)

Bivariate Normal Example

Let $f(\mathbf{x})$ be a bivariate normal density with means

$$\mu' = \begin{bmatrix} 1 & 2 \end{bmatrix}$$

and covariance matrix

$$\mathbf{V} = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2.0 \end{bmatrix}$$

Suppose we do not know how to draw samples from $f(\mathbf{x})$, but know how to draw samples from $f(x_i | x_j)$, which is univariate normal with mean:

$$\mu_{i.j} = \mu_i + \frac{v_{ij}}{v_{jj}}(x_j - \mu_j)$$

and variance

$$v_{i.j} = v_{ii} - \frac{v_{ij}^2}{v_{jj}}$$

```

In [125]: m = fill(0,2)
nSamples = 2000
m = [1.0, 2.0]
v = [1.0 0.5; 0.5 2.0]
y = fill(0.0,2)
sum = fill(0.0,2)
s12 = sqrt( v[1,1] - v[1,2]*v[1,2]/v[2,2])
s21 = sqrt(v[2,2] - v[1,2]*v[1,2]/v[1,1])
m1 = 0
m2 = 0;
for (iter in 1:nSamples)
    m12 = m[1] + v[1,2]/v[2,2]*(y[2] - m[2])
    m21 = m[2] + v[1,2]/v[1,1]*(y[1] - m[1])
    y[1] = rand(Normal(m12,s12),1)[1]
    y[2] = rand(Normal(m21,s21),1)[1]
    sum += y
    mean = sum/iter
    if iter%100 == 0
        @printf "%10d %8.2f %8.2f \n" iter mean[1] mean[2]
    end
end

```

100	1.09	2.21
200	1.06	2.16
300	1.06	2.16
400	1.05	2.12
500	1.03	2.11
600	1.01	2.10
700	1.00	2.09
800	1.01	2.09
900	1.00	2.08
1000	1.02	2.10
1100	1.00	2.09
1200	1.01	2.08
1300	1.01	2.08
1400	1.02	2.08
1500	1.03	2.10
1600	1.02	2.08
1700	1.02	2.08
1800	1.02	2.08
1900	1.03	2.07
2000	1.02	2.06

Metropolis-Hastings Algorithm

- Sometimes may not be able to draw samples directly from $f(x_i | \mathbf{x}_{i-})$
- Convergence of the Gibbs sampler may be too slow
- Metropolis-Hastings (MH) for sampling from $f(\mathbf{x})$:
- a candidate sample, y , is drawn from a proposal distribution $q(y|x^{t-1})$

$$x^t = \begin{cases} y & \text{with probability } \alpha \\ x^{t-1} & \text{with probability } 1 - \alpha \end{cases}$$

$$\alpha = \min\left(1, \frac{f(y)q(x^{t-1}|y)}{f(x^{t-1})q(y|x^{t-1})}\right)$$

- The samples from MH is a Markov chain with stationary distribution $f(x)$

Bivariate Normal Example

```

In [127]: nSamples = 10000
m = [1.0, 2.0]
v = [1.0 0.5; 0.5 2.0]
vi = inv(v)
y = fill(0.0,2)
sum = fill(0.0,2)

m1 = 0
m2 = 0
xx = 0
y1 = 0
delta = 1.0
min1 = -delta*sqrt(v[1,1])
max1 = +delta*sqrt(v[1,1])
min2 = -delta*sqrt(v[2,2])
max2 = +delta*sqrt(v[2,2])
z = y-m
denOld = exp(-0.5*z'*vi*z)
d1 = Uniform(min1,max1)
d2 = Uniform(min2,max2)
ynew = fill(0.0,2);
for (iter in 1:nSamples)
    ynew[1] = y[1] + rand(d1,1)[1]
    ynew[2] = y[2] + rand(d2,1)[1]
    denNew = exp(-0.5*(ynew-m)'*vi*(ynew-m));
    alpha = denNew/denOld;
    u = rand()
    if (u < alpha[1])
        y = copy(ynew)
        denOld = exp(-0.5*(y-m)'*vi*(y-m))
    end
    sum += y
    mean = sum/iter
    if iter%1000 == 0
        @printf "%10d %8.2f %8.2f \n" iter mean[1] mean[2]
    end
end
end

```

1000	1.04	1.93
2000	1.10	1.91
3000	1.13	1.91
4000	1.13	1.98
5000	1.05	1.96
6000	1.03	1.94
7000	1.03	1.96
8000	1.03	1.96
9000	1.04	1.96
10000	1.06	1.97

